

A tall, elegant glass filled with a golden beer, topped with a thick, white head of foam. The glass is set against a dark, blurred background, possibly a bar or restaurant setting.

NINKASI: Beer Recommender

Alex Li, Andrew Wu,
Nelson Chen, Luke Chu

Jan 17, 2016

A circular logo with a dark background and a thin white border. Inside the circle, the letters 'NK' are written in a bold, white, sans-serif font.

NK



Roadmap



Introduction

A tall, elegant glass filled with golden beer and a thick, white head of foam. The glass is condensation-covered and sits on a dark, reflective bar surface. The background is dark and out of focus.

NK



Motivation

- Ever get tired of drinking common beers like Budweiser or Corona? Move into craft beers?
- Want to find a similar beer?
- Want to find a different beer without taste testing too many before finding a good one?
- Let us recommend you some!





Competitor Research

- Only found systems to find similar beers
- Most sophisticated system, *Next Glass*, uses mass spectrometry to classify beers
- No recommenders that finds similar users





Our Goal:

Build hybrid beer recommendation system that finds both similar *items* and *users*





Project Workflow



Web Scraping

A tall, elegant glass filled with golden beer and a thick head of white foam sits on a metallic bar counter. The background is dark and blurred, suggesting a bar or restaurant setting.

NK



Data Scrapping



Home > Breweries > United States: Illinois > Goose Island Beer Company (AB-InBev)

Goose Island Bourbon County Stout



* picture credits
copyright may apply

overall
100
style
100

Brewed by **Goose Island Beer Company (AB-InBev)**

Style: **Imperial Stout** **Top 50**

Chicago, Illinois USA

Serve in Snifter

bottled common **on tap common** **Broad Distribution** **+**

[send corrections](#) | [shelftag](#) | [edit barcodes](#) | [update pic](#)

RATINGS: 2809 WEIGHTED AVG: 4.26/5 IBU: 60 EST. CALORIES: 426 ABV: 14.2%

COMMERCIAL DESCRIPTION

"I really wanted to do something special for our 1000th batch at the original brewpub. Goose Island could have thrown a party. But we did something better. We brewed a beer. A really big batch of stout-so big the malt was coming out of the top of the mash tun. After fermentation we brought in some bourbon barrels to age the stout. One hundred and fifty days later, Bourbon County Stout was born-A liquid as dark and dense as a black hole with a thick foam the color of bourbon barrels. The nose is a mix of charred oak, vanilla, caramel and smoke. One sip has more flavor than your average case of beer. It overpowers anything in the room. People have even said that it's a great cigar beer, but I haven't yet tried a cigar that would stand up to it." Brewmaster Greg Hall;
IBU's 60-High Color - Midnight

Was 11% abv,
2007 and 2008 - 13% abv
2011 - 14.5% abv
2012 - 15% abv
2013 - 14.9% abv
2014 - 14.4% abv
2015 - 14.2% abv

Editor's Note: Baudoinia Fulton & Wood Series offering (And the Low Storage entry from FoBAB 2012) is simply Bourbon County Stout aged in the same barrels. It has been altd as it offers no different recipe or barrel type simply the presence of a distillers fungus on the barrels that adds no distinct characteristic to the beer from the fungus itself except for possible oxygen exposure changes. While Brewers Intent indicates that they consider it a new beer, no true distinction aside from this oxidation amount and possible aging time differences separates the beers. It's essentially a single barrel Bourbon County version something we have always treated as regular Bourbon County Stout in the past.

[Most Recent](#) [Top Raters](#) [Highest Score](#) [Rated By](#) [Ticked By](#)

4 AROMA 9/10 APPEARANCE 3/5 TASTE 8/10 PALATE 4/5 OVERALL 16/20
Sudz4Dayz (70) **Montreal, Quebec, CANADA - DEC 16, 2016**

2016 bottle purchased at Fort Point Market in South Boston, MA. Pours jet black, smooth, cappuccino head with small frothy bubbles. Dissipates very quickly. Smells like burnt wood, vanilla, roastiness with bourbon sweetness and booze as well. Love it. Nice roasty flavor, smooth vanilla. Those come first and are quickly replaced by ample, everlasting dark chocolate flavor. Sweet bourbon barrel throughout and slight charred wood. Not the thickest mouthfeel but I wouldn't say light either. Medium. Overall not all too complex, balance is ok. The flavor is great and it definitely saves this one from being mediocre. In terms of overall quality it's definitely missing balance, a little complexity and a little on the mouthfeel. Delicious nonetheless. I enjoyed this one.

4.1 AROMA 8/10 APPEARANCE 3/5 TASTE 9/10 PALATE 4/5 OVERALL 17/20
chinchill (4448) - South Carolina, USA - DEC 15, 2016

12 oz bottle dated 29 August 2014 served in a Belgian snifter. Removing the cap indicated a secure seal, but this has very little carbonation. This is not ideal for the appearance, but the low carbonation seems about right for the smooth, full, high ABV body. Boozy and woody aroma with dark roasted grains and molasses. Sweet bourbon dominates the rich flavor. Moderately woody with faint vanilla and coffee in the finish. Mild bourbon barrels and dark bread in the aftertaste.

4.7 AROMA 9/10 APPEARANCE 5/5 TASTE 9/10 PALATE 5/5 OVERALL 19/20
Korcz (1142) - Warsaw, POLAND - DEC 11, 2016

Backlog, ocena przepisana z untappd, w ramach uzupełniania profilu na ratebeer. Genialne. Perfekcyjnie gładkie, oleiste, kremowe, eleganckie. Piękna wanilia, szlachetny alkohol, potężne nuty Bourbonu, odcieczkowy kokos i genialna jak na tę moc pijalność. Wybitne piwo, jeden z moich ulubionych stoutów ever. Jeśli nie ulubiony :)

4.2 AROMA 9/10 APPEARANCE 5/5 TASTE 8/10 PALATE 5/5 OVERALL 15/20
Listigovers (967) - Toronto, Ontario, CANADA - DEC 10, 2016

Pitch black, very strong whisky scent. Taste is very sweet and there is also heat from the 14%, tar like texture, very roasted and bourbony flavour. A true sipper, incredibly heavy and warming beer.



Data Set

- 1275 unique beers
- Top 25 beers per state
- 16,000+ users
- 280,000+ reviews/ratings




Exploratory Data Analysis

A tall, elegant glass filled with golden beer and a thick, white head of foam. The glass is condensation-covered and sits on a dark, reflective bar surface. The background is blurred, showing a dimly lit bar environment.

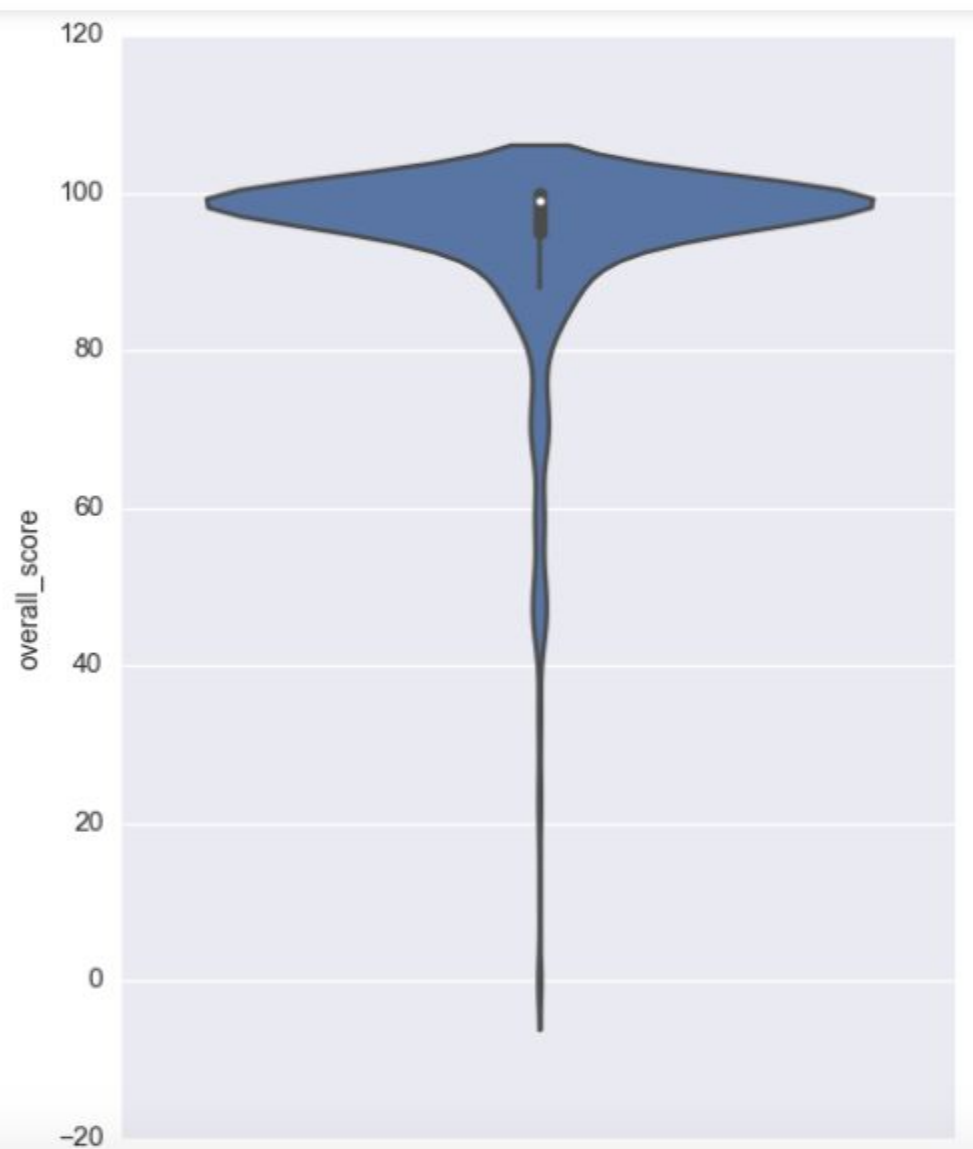
NK




Weighted Avg



Overall Score





| | count |
|----------------------|-------|
| Imperial Stout | 352 |
| Imperial IPA | 146 |
| India Pale Ale (IPA) | 118 |
| Sour/Wild Ale | 85 |
| Barley Wine | 54 |
| Saison | 45 |
| Mead | 42 |
| American Pale Ale | 42 |
| Porter | 40 |
| Imperial Porter | 35 |
| American Strong Ale | 30 |
| Sweet Stout | 26 |
| Black IPA | 26 |
| Stout | 24 |
| Belgian Strong Ale | 16 |
| Sour Red/Brown | 14 |
| Spice/Herb/Vegetable | 13 |
| Lambic Style - Fruit | 13 |
| Berliner Weisse | 12 |
| Baltic Porter | 11 |



Latent Dirichlet Allocation

Top Two Topics:

Stout

| Word | Proportion |
|-----------|------------|
| chocolate | 0.012 |
| coffee | 0.010 |
| bourbon | 0.009 |
| black | 0.006 |
| roasted | 0.006 |
| roast | 0.004 |
| brown | 0.004 |
| vanilla | 0.004 |
| dark | 0.003 |
| stout | 0.003 |

Fruity IPA

| Word | Proportion |
|------------|------------|
| pine | .006 |
| citrus | .006 |
| grapefruit | .006 |
| orange | .005 |
| ipa | .004 |
| tropical | .004 |
| hop | .004 |
| golden | .004 |
| funk | .003 |
| mango | .003 |

Recommender System



NK



Recommender System



Reviews

Text Processing

TF-IDF, LSI



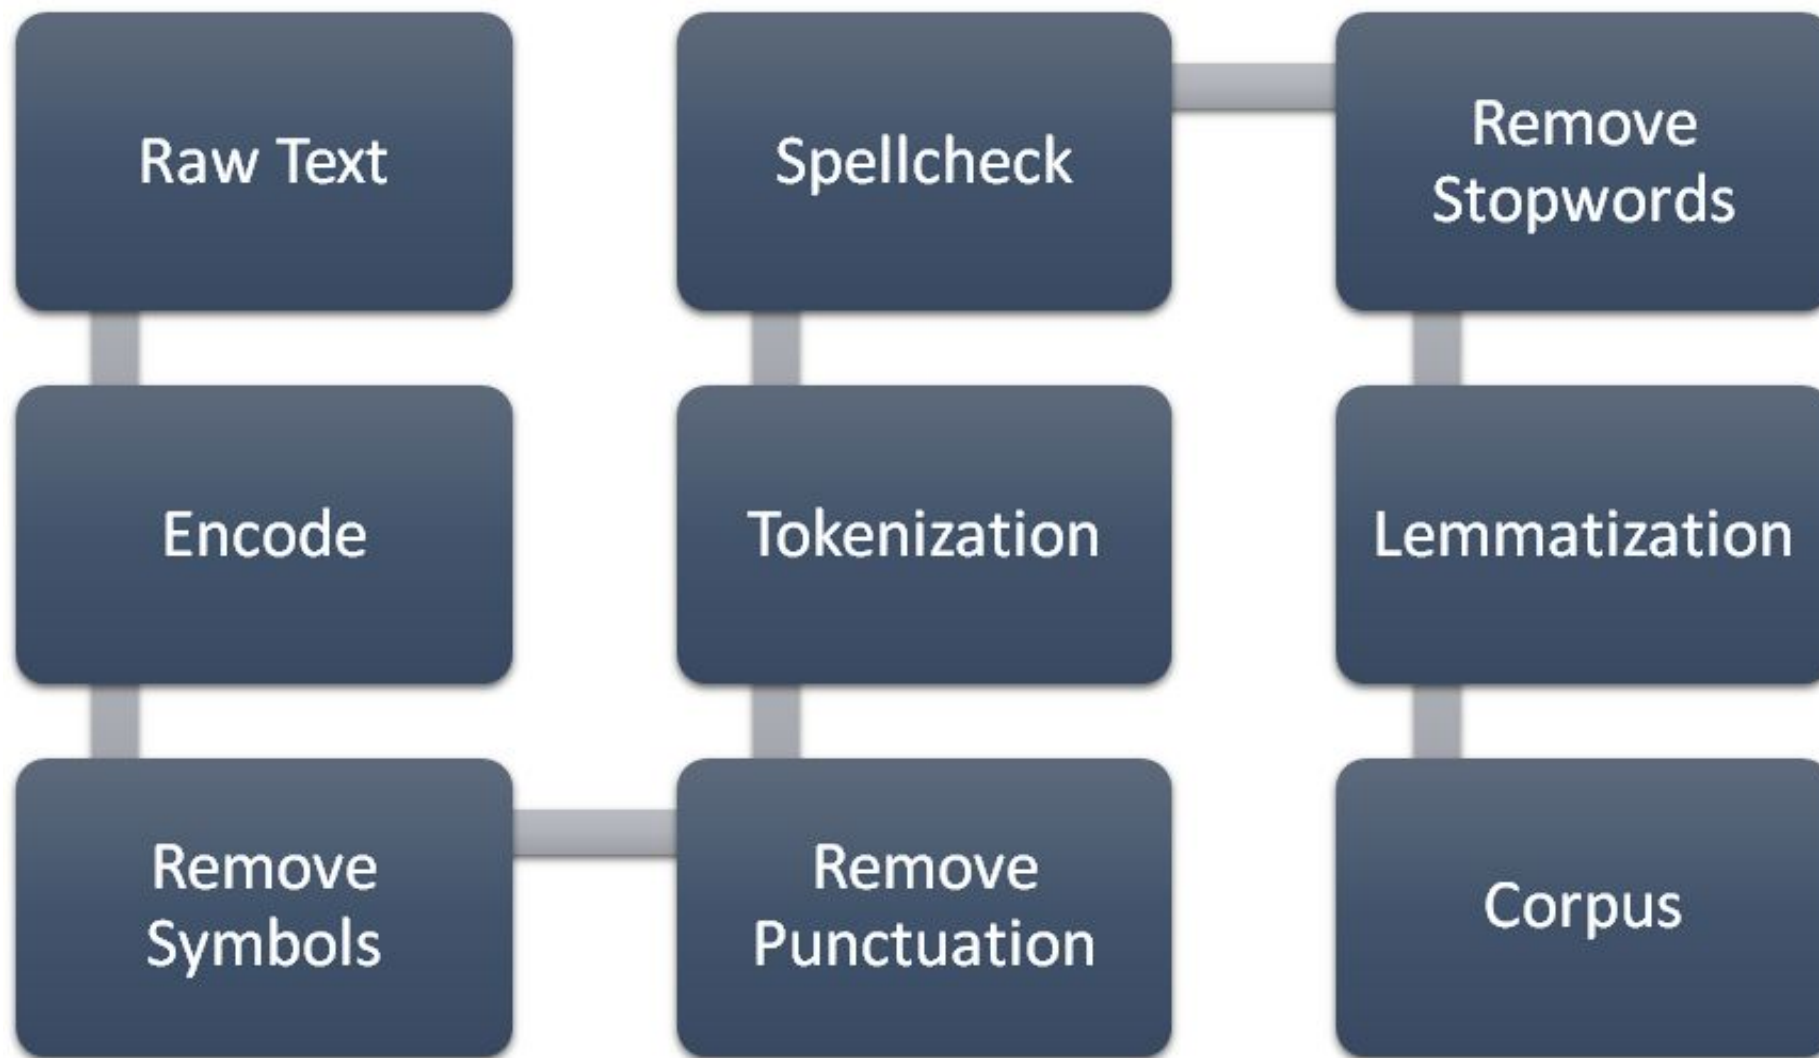
Ratings

SVD++, RBM

Ensemble



Content-Based: NLP





Content-Based: Algorithms

- Term Frequency-Inverse Document Frequency (TF-IDF): calculates importance of terms in each document
- Latent Semantic Indexing (LSI): does dimension reduction on the document-term matrix





Collaborative Filtering Algorithms

- Recommendation based on beer rating (explicit information)
- Two separate models were used
 - Singular Value Decomposition++ (SVD++)
 - Restricted Boltzmann Machine (RBM)
- Ensemble models together (Average)





Collaborative Filtering Results

- SVD++ : 91.4% Accuracy
- RBM : 90.7% Accuracy
- Ensemble: 92.1% Accuracy





Prediction Method for New Users

- Neighborhood approach using cosine similarity
- Impute missing ratings with weighted average
- Rank and output top 5 as recommendations



Flask App Demonstration



NK

ninkasibeer.herokuapp.com

Future Steps

A tall, elegant glass filled with golden beer and a thick, white head of foam. The glass is condensation-covered and sits on a dark, reflective bar surface. The background is dark and out of focus.

NK



Future Steps

- Expand beer database
- Add in more features (i.e style, palate, appearance, aroma, taste, and time)
- Build better single models (hyperparameter tuning, try new algorithms)
- Ensemble smarter (use minimizer or stacking)
- Develop App functionalities further





Thank you, Cheers!



Appendix


A tall, elegant glass filled with golden beer and a thick, white head of foam. The glass sits on a dark, reflective bar counter. The background is dark and out of focus, suggesting a bar or restaurant setting.

NK



SVD For Collaborative Filtering

- Latent Matrix Factorization of user-item matrix to find *latent* features
- Used Spark to implement as baseline: RMSE = 2.15





SVD++ (Implicit Feedback Version)

- Developed by Netflix Challenge winners
- Decomposes rating equation further
- Users that have less ratings are penalized more



$$\hat{r}_{ui} = \mu + b_u + b_i + q_i^T (p_u + \frac{1}{\sqrt{|N(u)|}} \sum_{j \in N(u)} y_j)$$

Diagram illustrating the SVD++ rating equation components:

- \hat{r}_{ui} : Predicted rating
- μ : Global average
- b_u : User bias
- b_i : Item bias
- q_i^T : Item Latent feature
- p_u : User latent feature
- $\frac{1}{\sqrt{|N(u)|}}$: Implicit feedback
- $\sum_{j \in N(u)} y_j$: Implicit feedback parameters



Restricted Boltzmann Machine (RBM)


- RBM is a generative, two-layer neural network
- Inspired by the Boltzmann Distribution from thermodynamics and fluid dynamics to minimize the energy function
- Performs CF by reconstructing the user-item matrix
- Hidden units can be thought of as latent features





Restricted Boltzmann Machine

- To train a RBM:
 - Initialize the visible layer
 - Hidden layer is calculated
 - Missing values are imputed
 - Weights are updated based with gradient descent
- The RBM is trained for every user, but the weights and bias are shared across all users





Ensemble CF models

- Averaged SVD++ and RBM predictions to get final predictions





Term Frequency-Inverse Document Frequency (TF-IDF)

- Calculate the “importance” of every word to a review in a corpus.
- Term Frequency is the number of times a word occurs in a document.

$$\text{tf}(t, d) = 0.5 + 0.5 \cdot \frac{f_{t,d}}{\max\{f_{t',d} : t' \in d\}}$$

- Inverse Document Frequency measures how rare is the word across all documents.

$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

- The TF-IDF is defined as

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$





Latent Semantic Analysis (LSA)

- LSA performs dimension reduction on the TF-IDF matrix

$$\begin{array}{ccccccc} & X & & U & & \Sigma & & V^T \\ & (\mathbf{d}_j) & & & & & & (\hat{\mathbf{d}}_j) \\ & \downarrow & & & & & & \downarrow \\ (\mathbf{t}_i^T) \rightarrow & \begin{bmatrix} x_{1,1} & \dots & x_{1,n} \\ \vdots & \ddots & \vdots \\ x_{m,1} & \dots & x_{m,n} \end{bmatrix} & = & (\hat{\mathbf{t}}_i^T) \rightarrow & \begin{bmatrix} \begin{bmatrix} \end{bmatrix} \\ \mathbf{u}_1 \end{bmatrix} \dots \begin{bmatrix} \end{bmatrix} \\ & & & & & & & \begin{bmatrix} \end{bmatrix} \\ & & & & & & & \mathbf{v}_1 \\ & & & & & & & \vdots \\ & & & & & & & \mathbf{v}_l \end{bmatrix} \end{array}$$

